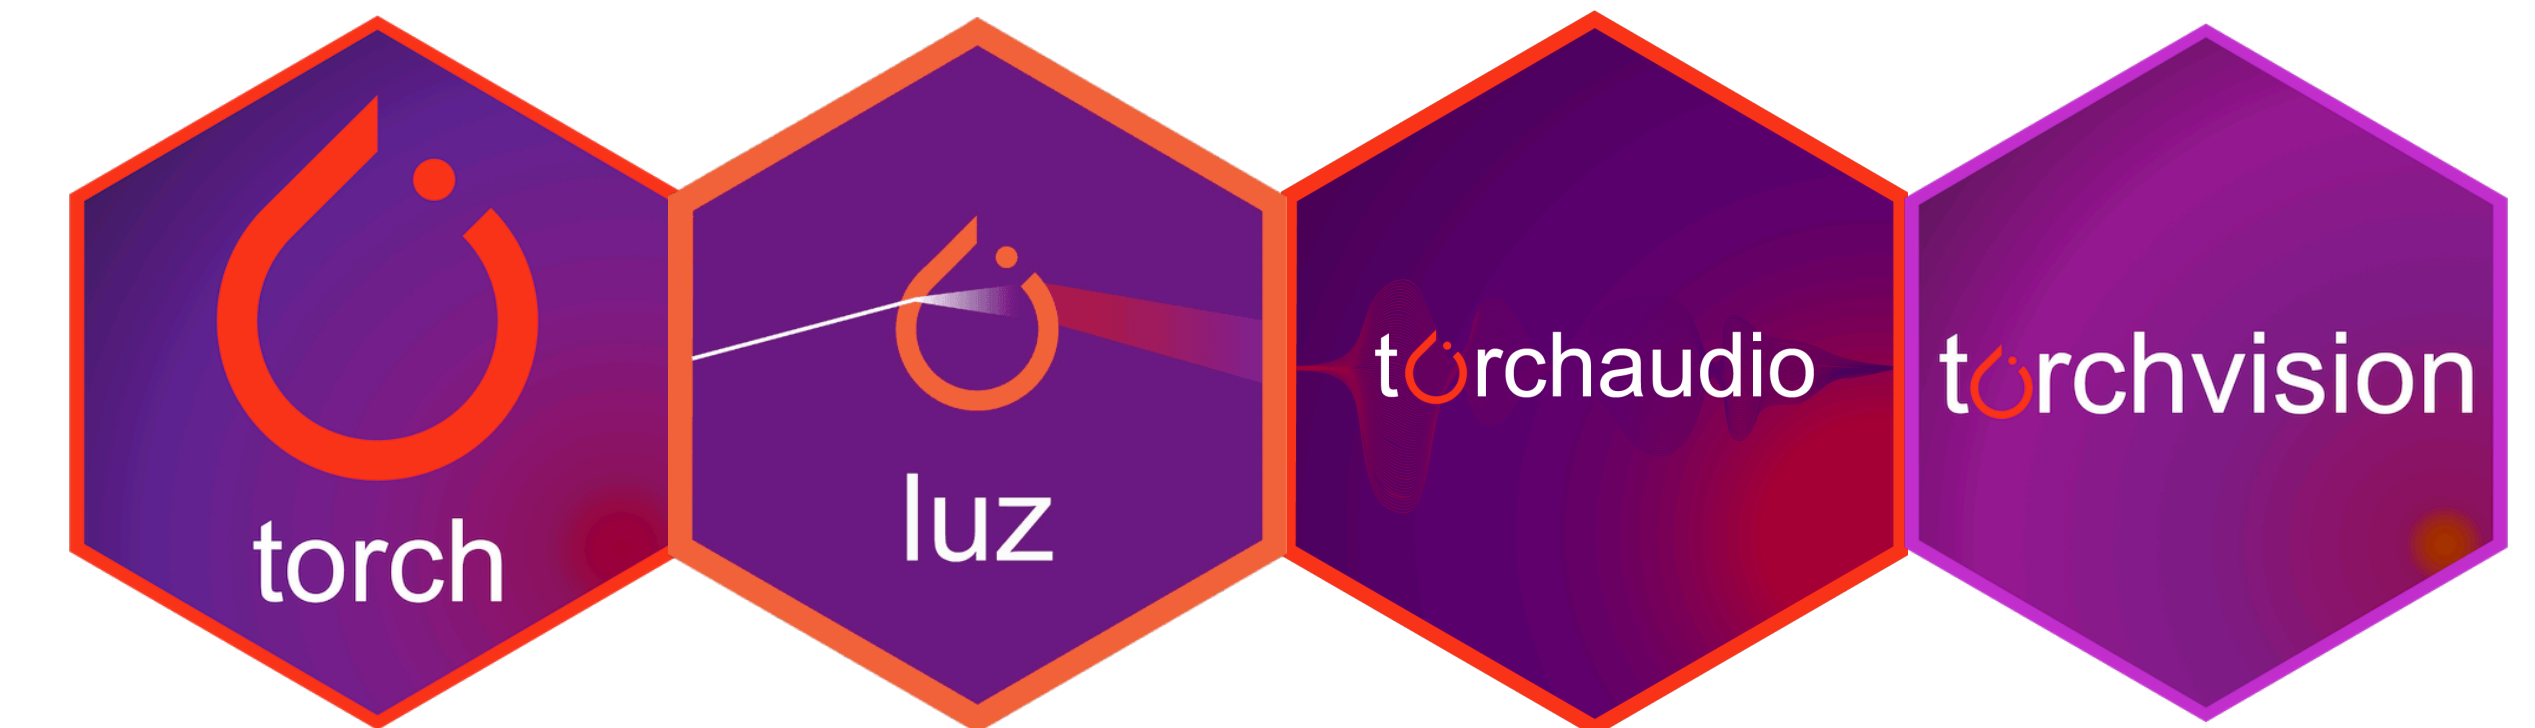


Deep Learning with {torch} CHEAT SHEET

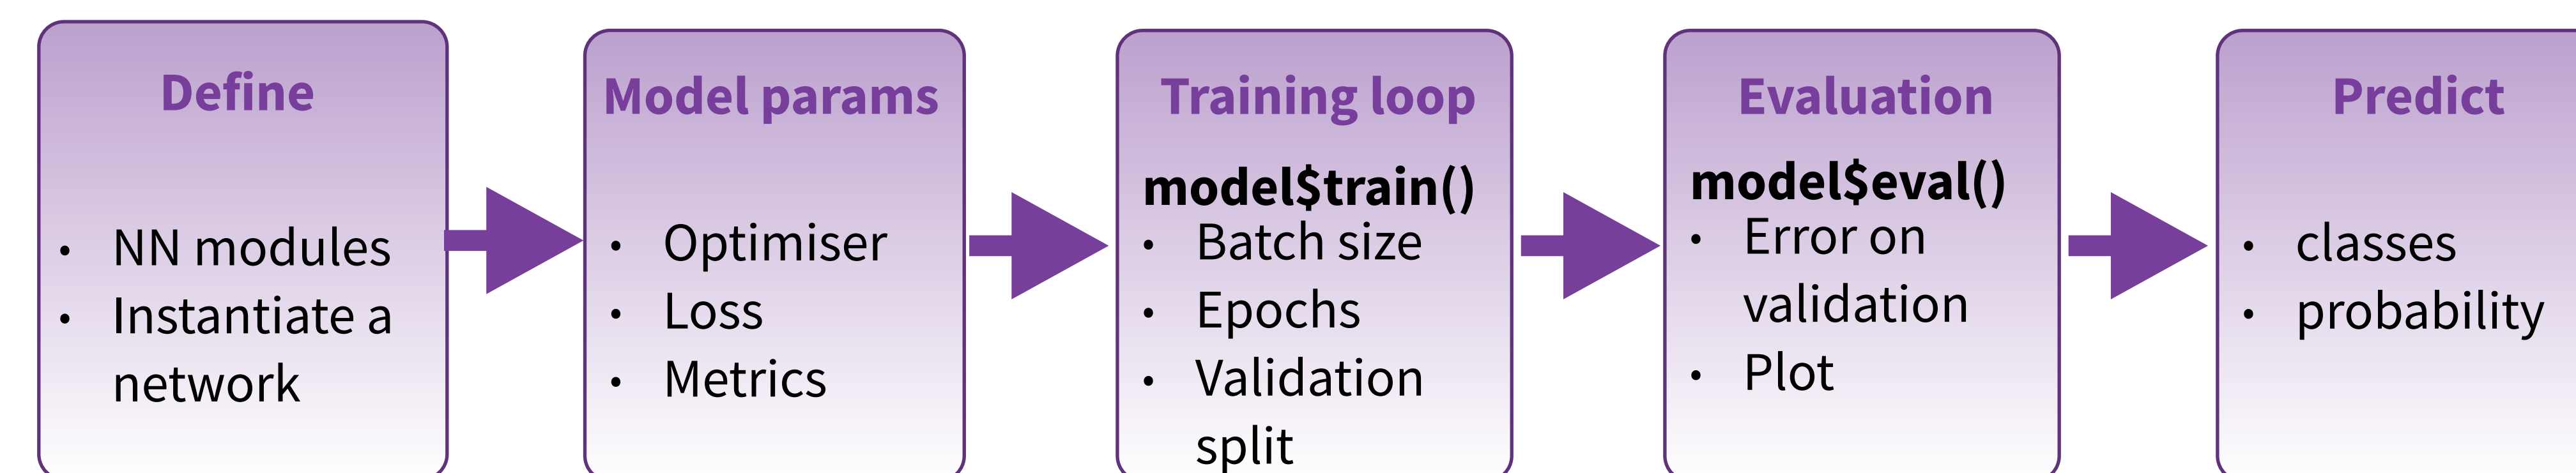


Intro

{torch} is based on Pytorch, a framework popular among deep learning researchers.

{torch}'s GPU acceleration allows to implement fast machine learning algorithms using its convenient interface, as well as a vast range of use cases, not only for deep learning, according to its flexibility and its low level API.

It is part of an ecosystem of packages to interface with specific dataset like {torchaudio} for timeseries-like, {torchvision} for image-like, and {tabnet} for tabular data. It is complemented by {lux} for a higher-level programming interface



<https://torch.mlverse.org/>

<https://mlverse.shinyapps.io/torch-tour/>

INSTALLATION

The torch R package uses the C++ libtorch library. You can install the prerequisites directly from R.

<https://torch.mlverse.org/docs/articles/installation.html>

```
install.packages("torch")
library(torch)
install_torch()
```

See ?install_torch for GPU instructions

Working with torch models

DEFINE A NN MODULE

```
dense <- nn_module(
  "no_biais_dense_layer",
  initialize = function(in_f, out_f) {
    self$w <- nn_parameter(torch_randn(in_f, out_f))
  },
  forward = function(x) {
    torch_mm(x, self$w)
  }
)
```

Create a nn module names no_biais_dense_layer

ASSEMBLE MODULES INTO NETWORK

```
model <- dense(4, 3)
```

Instantiate a network from a single module

```
model <- nn_sequential(
```

```
  dense(4,3), nn_relu(), nn_dropout(0.4),
```

```
  dense(3,1), nn_sigmoid())
```

Instantiate a sequential network with multiple layers

MODEL FIT

```
model$train()
```

Turns on gradient update

```
with_enable_grad({
  y_pred <- model(trainset)
  loss <- (y_pred - y)$pow(2)$mean()
  loss$backward()
})
```

Detailed training loop step (alternative)

EVALUATE A MODEL

```
model$eval()
```

or

```
with_no_grad({
  model(validationset)
})
```

Perform forward operation with no gradient update

OPTIMIZATION

```
optim_sgd()
```

Stochastic gradient descent optimiser

```
optim_adam()
```

ADAM optimiser

CLASSIFICATION LOSS FUNCTION

```
nn_cross_entropy_loss()
```

```
nn_bce_loss()
```

```
nn_bce_with_logits_loss()
```

(Binary) cross-entropy losses

```
nn_nll_loss()
```

Negative log-likelihood loss

```
nn_margin_ranking_loss()
```

```
nn_hinge_embedding_loss()
```

```
nn_multi_margin_loss()
```

```
nn_multilabel_margin_loss()
```

(Multiclass) (multi label) hinge losses

REGRESSION LOSS FUNCTION

```
nn_l1_loss()
```

L1 loss

```
nn_mse_loss()
```

MSE loss nn_ctc_loss()

Connectionist Temporal Classification loss

```
nn_cosine_embedding_loss()
```

Cosine embedding loss

```
nn_kl_div_loss()
```

Kullback-Leibler divergence loss

```
nn_poisson_nll_loss()
```

Poisson NLL loss

OTHER MODEL OPERATIONS

```
summary()
```

Print a summary of a torch model

```
torch_save(); torch_load()
```

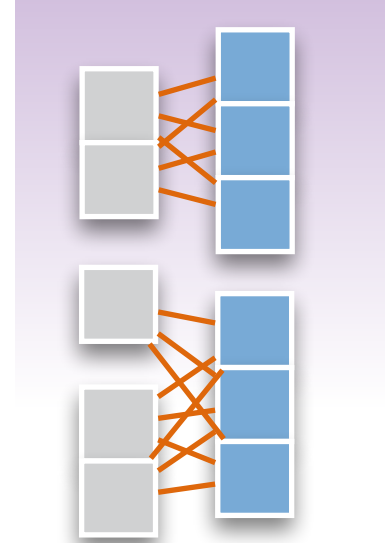
Save/Load models to files

```
load_state_dict()
```

Load a model saved in python

Neural-network layers

CORE LAYERS



```
nn_linear()
```

Add a linear transformation NN layer to an input

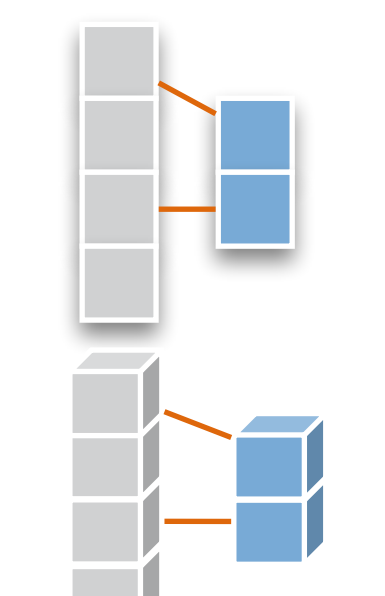
```
nn_bilinear()
```

to two inputs



```
nn_sigmoid(), nn_relu()
```

Apply an activation function to an output

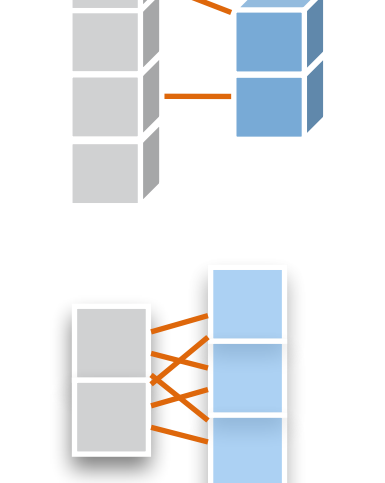


```
nn_dropout()
```

```
nn_dropout2d()
```

```
nn_dropout3d()
```

Applies Dropout to the input



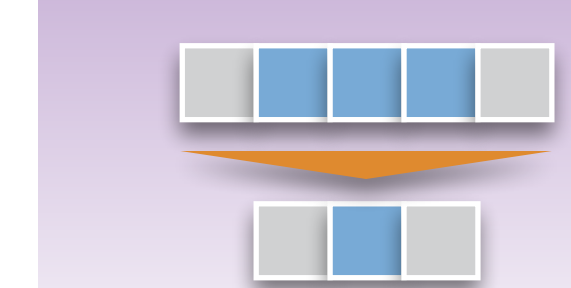
```
nn_batch_norm1d()
```

```
nn_batch_norm2d()
```

```
nn_batch_norm3d()
```

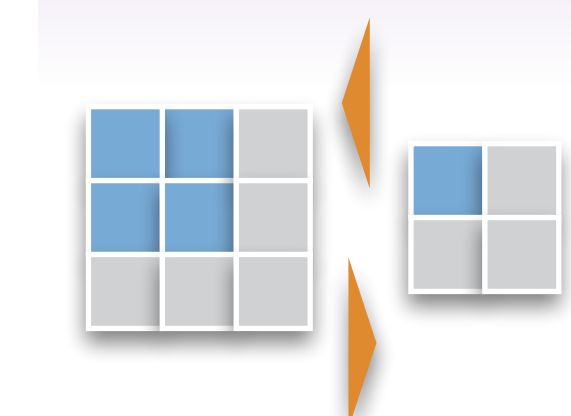
Applies batch normalisation to the weights

CONVOLUTIONAL LAYERS



```
nn_conv1d()
```

1D, e.g. temporal convolution



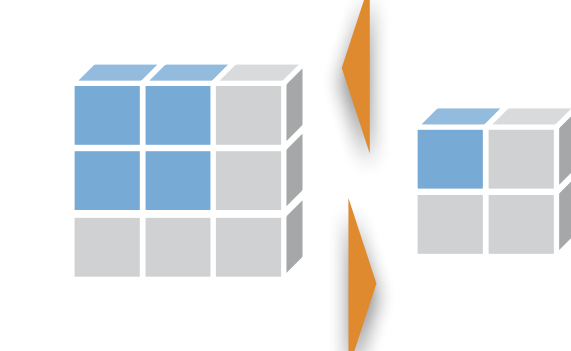
```
nn_conv_transpose2d()
```

Transposed 2D (deconvolution)



```
nn_conv2d()
```

2D, e.g. spatial convolution over images

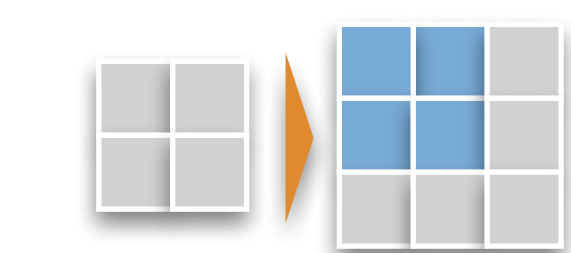


```
nn_conv_transpose3d()
```

Transposed 3D (deconvolution)

```
nn_conv3d()
```

3D, e.g. spatial convolution over volumes



```
nnf_pad()
```

Zero-padding layer

ACTIVATION LAYERS



```
nn_leaky_relu()
```

Leaky version of a rectified linear unit



```
nn_relu6()
```

rectified linear unit clamped by 6



```
nn_rrelu()
```

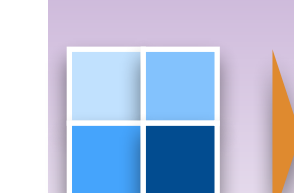
Randomized leaky rectified linear unit



```
nn_elu(), nn_selu()
```

Exponential linear unit, Scaled Exp lineal unit

POOLING LAYERS

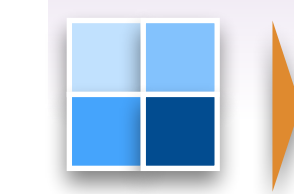


```
nn_max_pool1d()
```

```
nn_max_pool2d()
```

```
nn_max_pool3d()
```

Maximum pooling for 1D to 3D

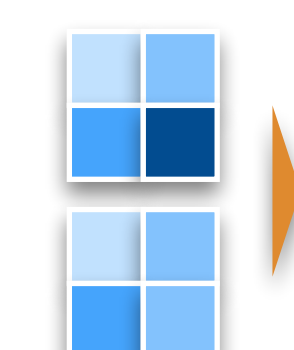


```
nn_avg_pool1d()
```

```
nn_avg_pool2d()
```

```
nn_avg_pool3d()
```

Average pooling for 1D to 3D

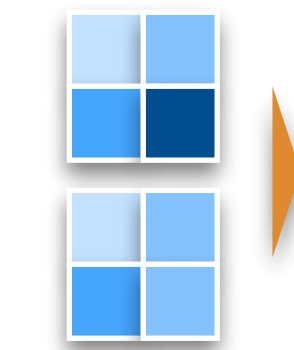


```
nn_adaptive_max_pool1d()
```

```
nn_adaptive_max_pool2d()
```

```
nn_adaptive_max_pool3d()
```

Adaptive maximum pooling



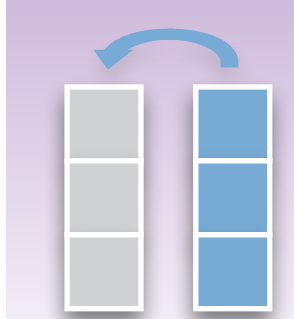
```
nn_adaptive_avg_pool1d()
```

```
nn_adaptive_avg_pool2d()
```

```
nn_adaptive_avg_pool3d()
```

Adaptive average pooling

RECURRENT LAYERS



```
nn_rnn()
```

Fully-connected RNN where the output is to be fed back to input

```
nn_gru()
```

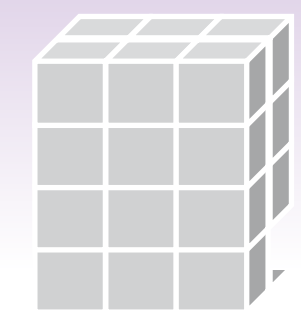
Gated recurrent unit - Cho et al

```
nn_lstm()
```

Long-Short Term Memory unit - Hochreiter 1997

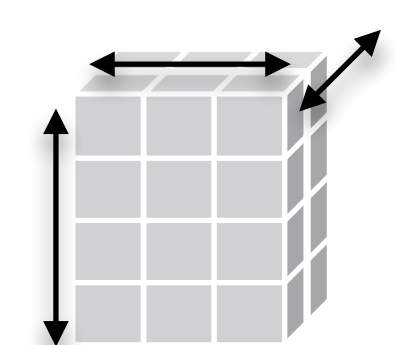
Tensor manipulation

TENSOR CREATION



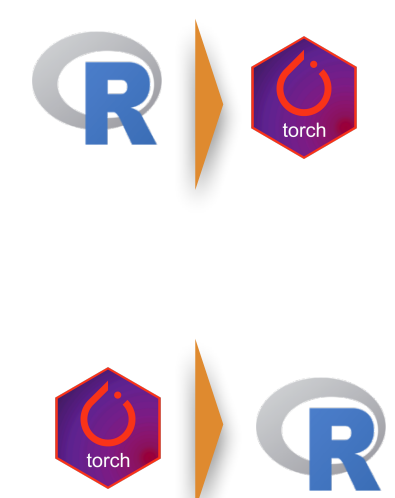
tt <- torch_rand(4,3,2) uniform distrib.
tt <- torch_randn(4,3,2) unit normal distrib.
tt <- torch_randint(1,7,c(4,3,2)) uniform integers within [1,7)
 Create a random values tensor with shape

tt <- torch_ones(4,3,2)
 torch_ones_like(a)
 Create a tensor full of 1 with given shape, or with the same shape as 'a'. Also **torch_zeros**, **torch_full**, **torch_arange**,...



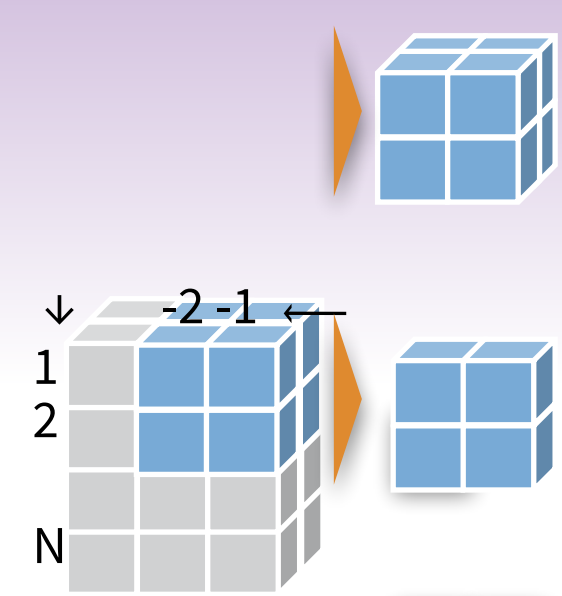
tt\$shape [1] 4 3 2 **tt\$ndim** [1] 3 **tt\$dtype** torch_Float
tt\$requires_grad [1] FALSE **tt\$device** torch_device(type='cpu')
 Get 't' tensor shape and attributes

tt\$stride() [1] 6 2 1
 jump needed to go from one element to the next in each dimension



tt <- torch_tensor(a, dtype=torch_float(), device="cuda")
 Copy the R array 'a' into a tensor of float on the GPU
a <- as.matrix(tt\$to(device="cpu"))

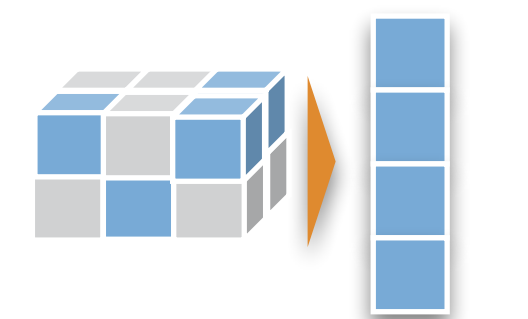
TENSOR SLICING



tt[1:2, -2:-1,]
 Slice a 3D tensor
tt[5:N, -2:-1, ..]
 Slice a 3D or more tensor, N for last

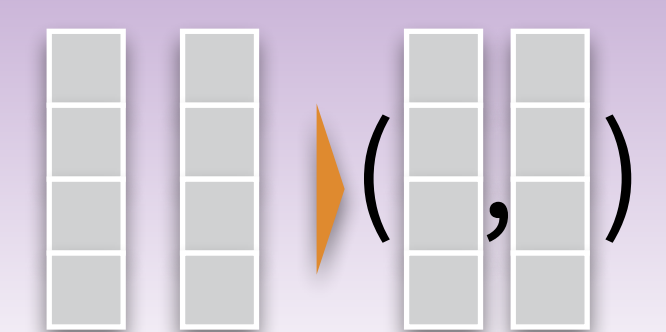
tt[1:2, -2:-1, 1:1]
tt[1:2, -2:-1, 1, keep=TRUE]
 Slice a 3D and keep the unitary dim.

tt[1:2, -2:-1, 1]
 Slice by default remove unitary dim.

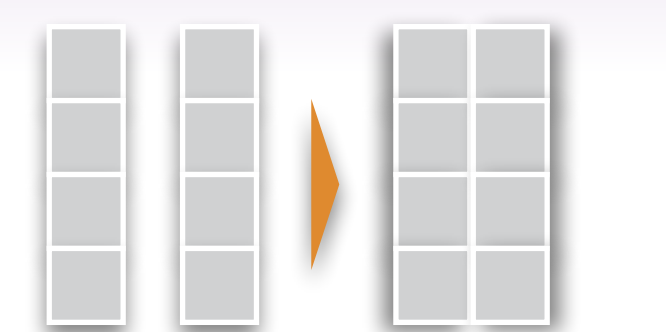


tt[tt > 3.1]
 Boolean filtering (flattened result)

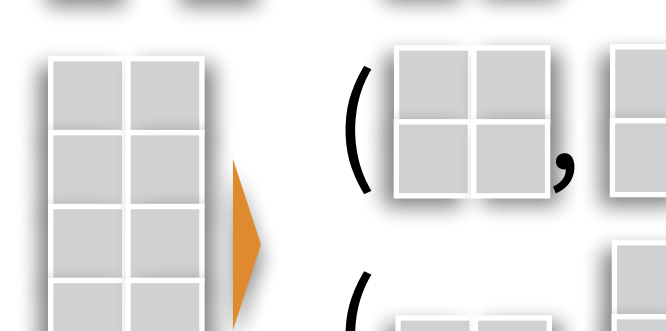
TENSOR CONCATENATION



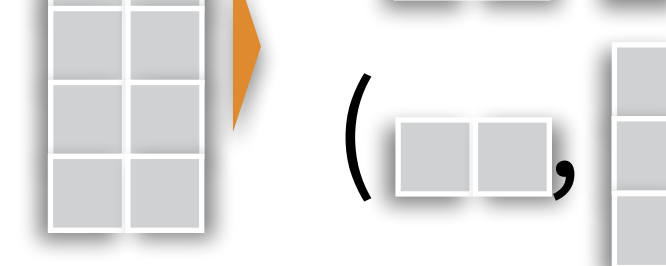
torch_stack()
 Stack of tensors



torch_cat()
 Assemble tensors

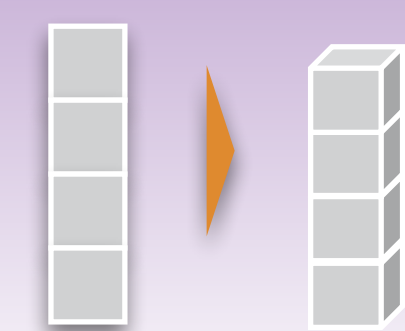


torch_split(2)
 split tensor in sections of size 2

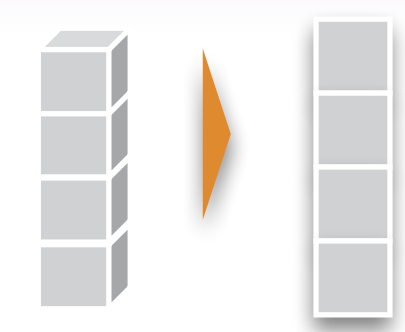


torch_split(c(1,3,1))
 split tensor into explicit sizes

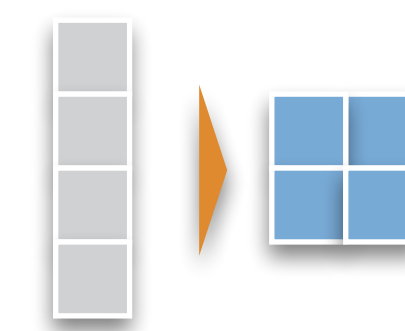
TENSOR SHAPE OPERATIONS



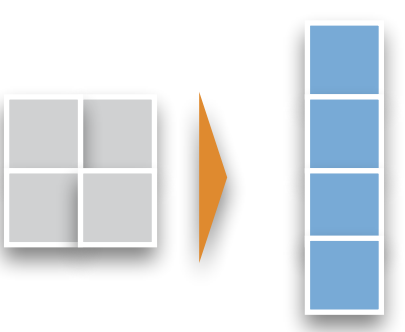
tt\$unsqueeze(1)
torch_unsqueeze(tt,1)
 Add a unitary dimension to tensor "tt" as first dimension



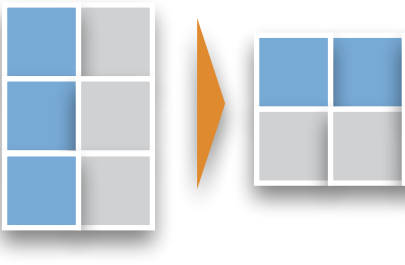
tt\$squeeze(1)
torch_squeeze(tt,1)
 Remove first unitary dimension to tensor "tt"



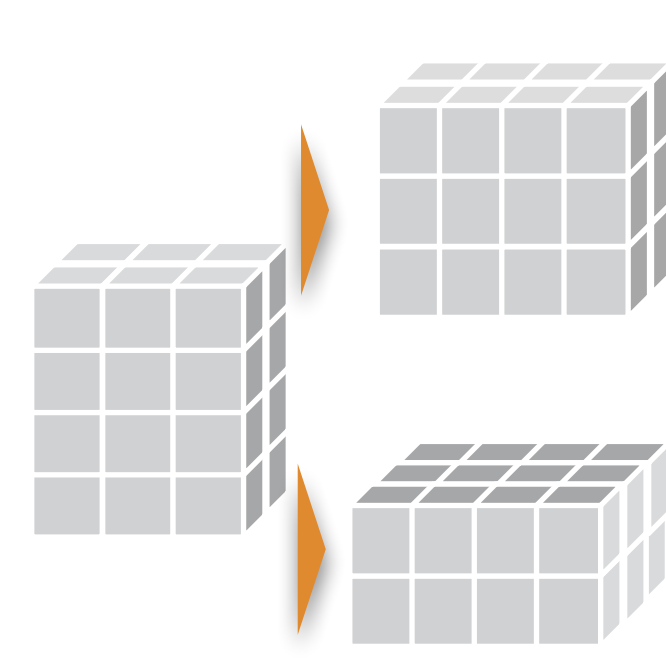
torch_reshape() \$view()
 Change the tensor shape, with copy or (tentatively) without



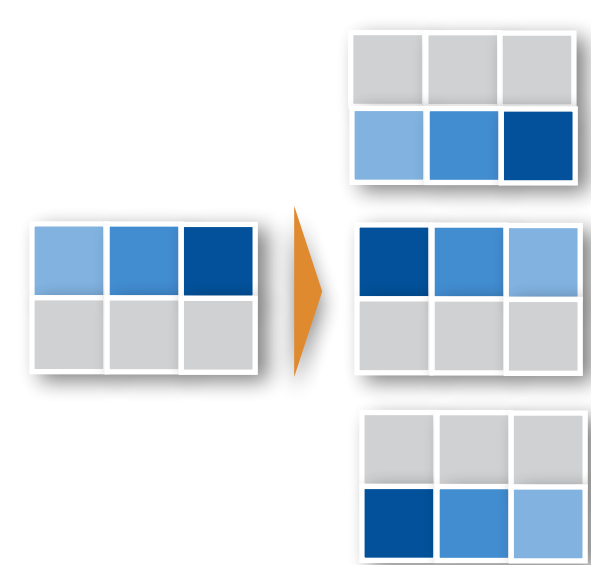
torch_flatten()
 Flattens an input



torch_transpose()

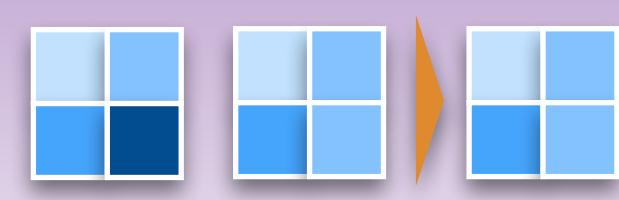


torch_movedim(c(1,2))
 switch dimension 1 with 2
torch_movedim(c(1,2,3), c(3,1,2))
 move dim 1 to dim 3, dim 2 to 1, dim 3 to 2
torch_permute(c(3,1,2))
 Only provide the target dimension order

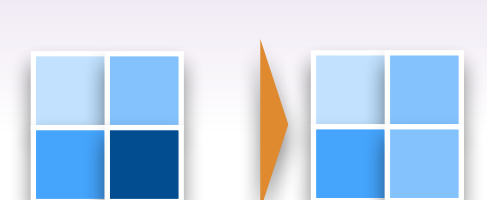


torch_flip(1) flip values along dim 1
torch_flip(2) 2
torch_flip(c(1,2)) both dims

TENSOR VALUES OPERATIONS



+, -, *
 Operations with two tensors

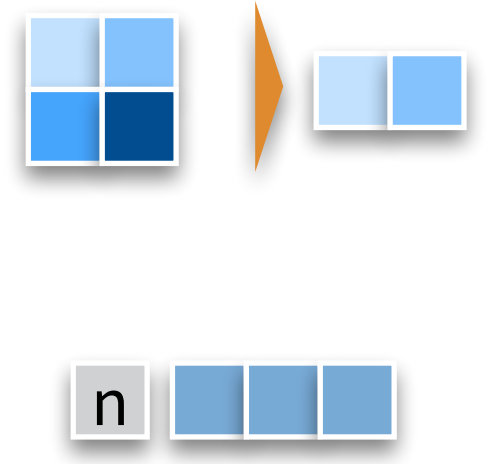


\$pow(2), \$log(), \$exp(), \$abs(), \$floor(), \$round(), \$cos(), \$fmod(3), \$fmax(1), \$fmin(3)
torch_clamp(tt, min=0.1, max=0.7)
 Element-wise operations on a tensor

\$eq(), \$ge(), \$le()
 Element-wise comparison

\$to(dtype = torch_long())
 Mutate values type

\$sum(dim=1), \$mean(), \$max()
\$amax()
 Aggregation functions on a single tensor



torch_repeat_interleave()
 Repeats the input n times



TRAINING AN IMAGE RECOGNIZER ON MNIST DATA

5041

The "Hello, World!" of deep learning

Pre-trained models

Torch applications are deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning.

NATIVE R MODELS

library(torchvision)
resnet34 <- model_resnet34(pretrained=TRUE)
 Resnet image classification model
resnet34_headless <- nn_prune_head(resnet34, 1)
 Remove top layer of a model

IMPORTING FROM PYTORCH

{torchvisionlib} allows you to import a pytorch model without recoding its nn modules in R. This is done in two steps

1- instantiate the model in Python, script it, and save it:

```
import torch
import torchvision
```

```
model = torchvision.models.segmentation.fcn_resnet50(pretrained = True)
model.eval()
```

```
scripted_model = torch.jit.script(model)
torch.jit.save(scripted_model, "fcn_resnet50.pt")
```

2- load and use the model in R:

```
library(torchvisionlib)
model <- torch::jit_load("fcn_resnet50.pt")
```

Troubleshooting

HELPERS

with_detect_anomaly()
 Provides insight of a nn_module() behaviour

Callbacks

A callback is a set of functions to be applied at given stages of the training procedure. You can use callbacks to get a view on internal states and statistics of the model during training.

```
# load MNIST images through a data loader
library(torchvision)
train_ds <- mnist_dataset( root = "~/cache",
  download = TRUE,
  transform = torchvision::transform_to_tensor
)
test_ds <- mnist_dataset( root = "~/cache",
  train = FALSE,
  transform = torchvision::transform_to_tensor
)
train_dl <- dataloader(train_ds, batch_size = 32,
  shuffle = TRUE)
test_dl <- dataloader(test_ds, batch_size = 32)

# defining the model and layers
net <- nn_module(
  "Net",
  initialize = function() {
    self$fc1 <- nn_linear(784, 128)
    self$fc2 <- nn_linear(128, 10)
  },
  forward = function(x) {
    x %>% torch_flatten(start_dim = 2) %>%
      self$fc1() %>% nnf_relu() %>%
      self$fc2() %>% nnf_log_softmax(dim = 1)
  }
)
model <- net()
# define optimizer
optimizer <- optim_sgd(model$parameters, lr = 0.01)
# train (fit)
for (epoch in 1:10) {
  train_losses <- c()
  test_losses <- c()
  for (b in enumerate(train_dl)) {
    optimizer$zero_grad()
    output <- model(b[[1]]$to(device = device))
    loss <- nnf_nll_loss(output, b[[2]]$to(device = device))
    loss$backward()
    optimizer$step()
    train_losses <- c(train_losses, loss$item())
  }
  for (b in enumerate(test_dl)) {
    model$eval()
    output <- model(b[[1]]$to(device = device))
    loss <- nnf_nll_loss(output, b[[2]]$to(device = device))
    test_losses <- c(test_losses, loss$item())
    model$train()
  }
}
```